

Regularized Synthetic Control Methods:

Advancing Causal Inference in Time Series Econometrics and
Observational Studies

Jörg Breitung^a, Lennart Bolwin^a, Justus Töns^a

^aUniversity of Cologne, Chair of Statistics and Econometrics

Abstract

The Synthetic Control (SC) method is a widely used tool for measuring causal treatment effects in observational trials. Typically, the counterfactual of the single treated unit is synthesized using a weighted average of the remaining units in the post-treatment phase. These weights are computed in a data-driven manner and aim to minimize the distance between the treated unit and its counterfactual in the pre-treatment phase. To avoid overfitting the training data and to ensure external validity of the results, the method's developers (Abadie, Diamond, and Hainmueller (ADH)) incorporated the constraint that each weight must be weakly positive, and all weights must sum up to one. Building on the work of [Doudchenko and Imbens, 2016], we propose a generalization that allows for the inclusion of a constant term and negative weights. However, we develop the Regularized Synthetic Controls (REGSC) estimator, an alternative regularization approach that shrinks individual coefficients towards zero and the sum of coefficients towards one. Besides the crucial advantage of a closed form expression, Monte Carlo studies confirm that this regularization method dominates other estimators in data-generating processes with factor structure as was proposed by the inventors of SC. Next, we extend the approach to dynamic contexts and propose a regularized Autoregressive Distributed Lag (ARDL) model for optimal estimation of the counterfactual in time series settings. Again, simulations confirm the new method's potential to enhance the accuracy and robustness of causal effect estimation in time series econometrics and observational studies.

Keywords: *Synthetic Control; Observational Studies; Causal Inference; Regularization, Autoregressive Distributed Lag Models*

List of Acronyms

ADH Abadie, Diamond, and Hainmueller

ARDL Autoregressive Distributed Lag

GDP Gross Domestic Product

DGP Data Generating Process

iid independent and identically distributed

MSFE Mean Squared Forecast Error

MSPE Mean Squared Prediction Error

OLS Ordinary Least Squares

PC Principal Components

REGSC Regularized Synthetic Controls

RSS Residual Sum of Squares

SC Synthetic Control

USA United States of America

VAR Vector Autoregression

VARSC Vector Autoregressive Synthetic Control

Contents

1	Introduction	6
2	Literature Review 2-3 pages	7
2.1	Synthetic Control	7
2.2	Overview	7
2.3	Application	8
2.4	Methodological Background	8
2.5	Extensions/ Developments	8
2.6	Testing	9
2.7	Time Series Econometrics	9
3	Theory	10
3.1	ADH Case	11
3.2	Simple Static Extension	16
3.3	General Static Extension	18
3.4	General Dynamic Extension	20
4	Simulation	21
4.1	Static Data Generating Processes	21
4.2	Dynamic Data Generating Processes	25
5	Simulation	26
5.1	Static Data Generating Processes	26
5.2	Weakly Dynamic Data Generating Processes	26
5.3	Dynamic Data Generating Processes	26
5.3.1	Set up	26

<i>CONTENTS</i>	4
5.3.2 Employed Models	27
5.3.3 Results	28
6 Applications	29
6.1 The Economic Costs of Conflict	29
6.2 Estimating the Effect of California's Tobacco Control Program	30
6.3 The Economic Cost of the 1990 German Reunification	31
7 Conclusion	32
8 Appendix	38
8.1 Simple Static Extension	38
8.1.1 OLS Solution	38
8.1.2 SC Solution	38
8.1.3 Variances	38
8.2 Simulation Study	39
8.2.1 Static Simulation results	39

List of Figures

1	Region Exclusion Procedure of ADH	15
2	Example Factor-Data Generating Process (DGP)	23
3	GDP per Capita for the (synthetic) Basque	29
4	Cigarette Sales per Capita for the (synthetic) California	30
5	GDP per Capita for the (synthetic) West Germany	31

1. Introduction

This is work in progress

SC method is combined with Vector Autoregression (VAR). Method was introduced by ADH. Paper strcuture could be similar to the one by [Doudchenko and Imbens, 2016]

2. Literature Review 2-3 pages

Literature has been clustered. This is work in progress

2.1. Synthetic Control

The **SC** method was developed by Alberto Abadie and colleagues in a series of influential papers ([Abadie and Gardeazabal, 2003], [Abadie et al., 2010], [Abadie et al., 2015]). The method is designed to estimate the causal effect of a treatment in a setting with a single treatment unit and a number of potential control units. Pre- and post-treatment data are observed for the treatment and control units for the outcome of interest as well as for a set of covariates. The **SC**-procedure combines aspects of the matching and difference-in-difference literature and can therefore be interpreted as a relative of the causal inference literature introduced by [Rubin, 1974]. Similar to many other microeconomic methods, the objective is to distinguish causation from correlation and to assess the magnitude and significance of treatments in observational case studies.

In their canonical 2003 article, Abadie and Gardeazabal evaluate the causal economic effects of conflict using terrorist conflicts in the Basque Country as a comparative case study. In their specific application example, they find that terrorist conflicts caused the per capita Gross Domestic Product (**GDP**) of the treatment unit (Basque Country) to decline by about 10% relative to the synthesized control unit.

some more words on other findings and things they did The next appropriate setting for an application of the **SC** method was the introduction of a large-scale tobacco control program implemented in the state of California in the United States of America (**USA**) in 1988.

2.2. Overview

[Abadie, 2021] read.

[Athey and Imbens, 2016] read.

2.3. Application

[Born et al., 2019] read.

[Cho, 2020] read.

[Cunningham, 2021] read.

[Funke et al., 2020] read.

2.4. Methodological Background

[Hainmueller et al., 2011] read.

[Abadie and Imbens, 2006] not read.

[Abadie and Imbens, 2002] not read.

[Doudchenko and Imbens, 2016] read.

[Ferman, 2021] read.

[Frangakis and Rubin, 2002] not read.

[Rosenbaum and Rubin, 1983] not read

[Rubin, 1974] not read.

2.5. Extensions/ Developments

[Abadie and L'Hour, 2021] read.

[Amjad et al., 2018] read.

[Ben-Michael et al., 2021] read.

[Ben-Michael et al., 2021] not read.

[Kellogg et al., 2021] not read.

[Kuosmanen et al., 2021] not read.

[Muhlbach and Nielsen, 2019] read.

Developments

[Arkhangelsky et al., 2021] not read

[Athey et al., 2017] not read.

[Brodersen et al., 2015] read.

[von Brzeski et al., 2015] read.

[Hartford et al., 2017] read.

2.6. Testing

[[Andrews, 2003](#)] not read.

[[Cattaneo et al., 2021](#)] not read.

[[Chernozhukov et al., 2019](#)] not read.

[[Chernozhukov et al., 2021](#)] not read.

[[Firpo and Possebom, 2018](#)] not read.

[[Hahn and Shi, 2017](#)] read.

2.7. Time Series Econometrics

[[Martin et al., 2012](#)] read.

[[Harvey and Thiele, 2020](#)] read.

[[Breitung and Knüppel, 2021](#)] partially read.

3. Theory

In this chapter, we propose alternative **SC**-estimators to assess the magnitude of treatment effects in observational settings. To establish a general basis, we first describe the contextual environment of the estimation. Similar to the setting as introduced by **ADH**, we consider a framework with $J + 1$ panel units indexed by $j = 0, 1, \dots, J$ that are observed over a time horizon of T periods. Without loss of generality, assume that unit $j = 0$ is exposed to the treatment at period $t = T_0$ with $1 < T_0 < T$ and that there are no treatment anticipation and contamination (i.e., no spillovers in time and space). The former would be the case if the treatment affects unit $j = 0$ before T_0 , the latter describes the case where some of the supposedly untreated units $j = 1, \dots, J$ are contaminated as they are affected by the treatment. To contextualize these assumptions, [Abadie et al., 2010] argue that in the presence of anticipation effects, T_0 could be shifted back in time until the no-anticipation assumption seem plausible. If panel units in the donor pool¹ are affected by the treatment (contamination) as it is likely in the Brexit-application, those units could be removed from the sample prior to the estimation. Our goal is to evaluate the causal effect of the treatment, the specific functional form of which remains unspecified though. This is possible because the main goal of the **SC**-estimation lies in the precise estimation of the counterfactual. Since the treatment scenario is empirically observable, it is not necessary to specify the specific functional form of the it.

The following chapter is structured as follows: We first describe the canonical estimation procedure as proposed by **ADH**. Furthermore, **ADH** propose a model-invariant hypothesis testing approach. As this approach is employed in the further analysis, we also give a brief overview of these principles. Next we build intuition by considering a simple static scenario with only two donor units and one treatment unit. This setting is subsequently generalized to the case with more potential donors. Our extensions diverge from the setting of **ADH** in two key aspects: First, we remove the weight constraints, leading us to explore regularization as a means to prevent overfitting. Second, we analyze a situation without covariates which drastically reduces the data requirements and causes our algorithm to estimate the counterfactual with a significantly smaller information set.

¹ To ensure direct comparability with the **SC** literature, we adopt most of the commonly used terms. For example, control group units are labeled as 'donors'.

However, this fact leads us to the necessity of utilizing all available information in an efficient manner and establishes our main contribution: The integration of multivariate time series approaches into the **SC**-algorithm. The theoretical derivation of this estimator completes the chapter.

3.1. ADH Case

We start by presenting the **SC**-method and the testing procedure as introduced by **ADH**. For the sake of comparability and due to its notational clarity, we borrow the employed notation of Abadie and colleagues. In terms of structural design, we build on the thorough presentation of the **SC**-method and the related hypothesis testing procedure by [Firpo and Possebom, 2018].

Setup

The estimation task can be constituted by the potential outcome framework as introduced by [Neyman, 1923] and elaborated by [Rubin, 1974]. Let $y_{j,t}^I$ be the (potential) outcome for unit j at point t in the presence of the intervention. Likewise, let $y_{j,t}^N$ be the (potential) outcome for j at point t in the absence of the intervention. **ADH** define the treatment effect of the intervention as

$$\delta_{j,t} = y_{j,t}^I - y_{j,t}^N$$

and introduce the indicator variable $D_{j,t}$ that takes on the value 1 if unit j is treated at period t and the value 0 otherwise. Given the assumed absence of anticipation and contamination, the following outcome is observed

$$y_{j,t} + D_{j,t}\delta_{j,t} = \begin{cases} y_{j,t}^N & \text{(if } j = 0 \text{ and } t < T_0) \text{ or } j \geq 1, \\ y_{j,t}^N + \delta_{j,t} & \text{if } j = 0 \text{ and } t \geq T_0. \end{cases}$$

The goal to estimate the causal treatment effect $(\delta_{0,T_0}, \dots, \delta_{0,T})$ therefore boils down to the estimation of the counterfactuals of unit $j = 0$ in the post-treatment phase $(y_{0,T_0}, \dots, y_{0,T})$, i.e. on what trajectory would unit $j = 0$ have been, was there no intervention. The basic idea of **ADH** is to estimate these counterfactuals as a weighted average of the donor outcomes using a data-driven approach to compute the weights. Intuitively, the weights

are computed such that they optimally predict the outcomes and a set of time-invariant explanatory variables for the treatment unit in the pre-intervention phase, conditional on having a percentage interpretation. Thus, for the computation of the weights, we focus exclusively on the pre-intervention time periods $t \in \{1, 2, \dots, T_0 - 1\}$. Subsequently, the counterfactuals are extrapolated by applying the calculated weights to the post-intervention time periods $t \in \{T_0, T_0 + 1, \dots, T\}$.

Let $Y_j = (y_{j,1}, \dots, y_{j,T_0-1})'$ be the vector of observed pre-intervention outcomes for unit j .² To distinguish treatment unit and donors, **ADH** collect the treatment unit in the $((T_0 - 1) \times 1)$ -vector Y_0 and row-bind all donor unit vectors into the $((T_0 - 1) \times J)$ -matrix Y_1 . Moreover, a set of K time-invariant covariates of Y_j is observed for all panel units.³ Therefore, let X_0 denote the $(K \times 1)$ -vector of covariates for Y_0 and let X_1 denote the $(K \times J)$ -matrix of explanatory variables for Y_1 . To estimate the causal effect of the treatment, the **SC**-estimator estimates the counterfactuals $(\hat{y}_{0,1}, \dots, \hat{y}_{0,T_0}, \dots, \hat{y}_{0,T})$ of the single treated unit for the pre- and post-intervention phase as

$$\hat{y}_{0,t} = \sum_{j=1}^J \hat{w}_j y_{j,t}^N \quad \forall t \in \{1, \dots, T\}$$

The weights $(\hat{w}_1, \dots, \hat{w}_J)$ are constraint such that $\hat{w}_j \geq 0 \quad \forall j$ and $\sum_{j=1}^J \hat{w}_j = 1$. It is worth noting that this constraint requires the counterfactuals to belong to the convex hull of the donors as otherwise, \hat{Y}_0 will never match its true counterpart. [Abadie et al., 2010] argue that "the magnitude of discrepancy" should be calculated in advance of each **SC**-application. If the researcher finds that the pre-intervention values of Y_0 fall outside the convex hull of the donors, the usage of **SC** is not recommended. Formally, $(\hat{w}_1, \dots, \hat{w}_J)$ is the solution of the following nested optimization problem:

$$\hat{w}(v) = \arg \min_w \sum_{k=1}^K v_k \left(x_{0,k} - \sum_{j=1}^J w_j x_{j,k} \right)^2$$

with v being an arbitrary positive definite vector of dimension $(K \times 1)$ which solve the

² For instance, in the canonical example of [Abadie and Gardeazabal, 2003], Y_0 would be the vector of **GDPs** for Great Britain until the Brexit referendum.

³ In the already mentioned Brexit-example, natural predictors of **GDP** are its components consumption, investment, government spending and net exports.

second optimization problem:

$$\hat{v} = \arg \min_v \sum_{t=1}^{T_0-1} \left(y_{0,t} - \sum_{j=1}^J \hat{w}_j(v) y_{j,t} \right)^2$$

Afterwards, the causal effect of the intervention $\delta_{j,t}$ can be quantified at each time point after the intervention $t \in \{T_0, T_0 + 1, \dots, T_1\}$ as the gap between observed ($y_{0,t}^N + \delta_{j,t}$) and predicted outcome ($\hat{y}_{0,t}^N$).

This two-step estimation procedure serves two crucial purposes: \hat{v} measures the relative importance of the K variables in X_1 to explain X_0 . In contrast, the weighting vector $\hat{w}(v)$ quantifies the relative importance of each unit in the donor pool. Summarizing the key concept of **ADH**, the **SC**-method ensures that the synthesized treatment unit is as similar as possible to the actual treatment unit with respect to the quantity of interest and a set of potential explanatory variables in the pre-treatment period. Especially in the canonical examples of **SC**, the quantity of interest (e.g. **GDP**) and the explanatory variables (e.g. consumption, investment, government spending and net exports) are interconnected by construction. Thus, observing that the **SC**-estimator was capable of approximating both targets significantly enhanced the methods credibility. If the explanatory variables are omitted, the **SC**-algorithm reduces to an Ordinary Least Squares (**OLS**) estimation, constraint to have no constant and weakly positive coefficients that sum up to one.

Hypothesis Testing

The question of treatment effect significance arises naturally subsequent to the construction of the synthetic control. **ADH** propose a model-invariant non-parametric inference procedure that is based on the Exact Hypothesis Test proposed by [Fisher, 1935]. The basic idea behind such permutation tests is to compare the observed data with a number of randomly permuted versions of it, and to use the distribution of the test statistic calculated of the permuted samples to estimate the probability that the observed result occurred by chance alone.

In the context of **SC**, **ADH** consider permutations in region (i.e. panel unit) and time. Region permutations estimate the treatment effect vector $(\delta_{j,T_0}, \dots, \delta_{j,T})$ for each panel

unit $j \in \{0, \dots, J\}$.⁴ This procedure provides the researcher with the empirical $(J + 1)$ -observational distribution of the treatment. Next, it is possible to compare the estimated treatment vector $(\delta_{0,T_0}, \dots, \delta_{0,T})$ of the truly treated unit with the J placebo-treatment vectors of the units of the donor pool. Given the estimated treatment effect for $j = 0$ is large, the null hypothesis of no treatment effect can be rejected at the significance level of one minus the percentile of $(\delta_{0,T_0}, \dots, \delta_{0,T})$ in the empirical distribution.⁵ Time permutations on the other hand consider only panel unit $j = 0$, permute T_0 to dates prior to the true treatment date and compute again the empirical treatment distribution. Given that $T_0 \gg J$, this approach can increase the sensitivity of the test, since the theoretically feasible significance threshold of region permutation tests is determined by $\frac{1}{J}$. For both, region and time permutations, **ADH** condense the vector of estimated treatment effects into a precision metric like the Mean Squared Forecast Error (**MSFE**)⁶ of the following form:

$$MSFE_j = \frac{\sum_{t=T_0}^T (\hat{y}_{j,t}^N - y_{j,t}^N)^2}{T - T_0}$$

A possible problem that can occur when assessing the relative rarity of the estimated treatment effect using the procedure described above is the existence of outliers in the donor pool. In the context of region permutations, suppose that a donor region is very different from the rest such that it falls outside the convex hull of the remaining donors. Note, that this circumstance does not cause problems for the truly treated region and its synthesized counterfactual as we expect the **SC**-algorithm to assign a near zero weight to such an outlier. However, since the outlier itself cannot be synthesized precisely by the donor pool, both **MSPE** and **MSFE** are expected to be large. As this special feature causes the permutation test to be unreasonably conservative, **ADH** propose to exclude regions that are hard to predict, i.e. who have a **MSPE** that exceeds the **MSPE** of the truly treated unit to a great extent. Figure 1 visualizes the exclusion procedure in the tobacco control application of **ADH**.

⁴ Note that it is necessary to exclude the truly treated unit from donor pool to ensure the validity of the no contamination assumption.

⁵ For instance, let $J = 99$ such that treatment effects for 100 panel units can be computed. As long as the estimated treatment effect of the truly treated units belongs to the 95 largest effects (95th percentile or higher), the permutation test rejects the null hypothesis of no treatment effect at least at 5 percent.

⁶ Note that **ADH** speak of the Mean Squared Prediction Error for dates before and after T_0 . Since we consider the time span until T_0 as prediction window and the time span after T_0 as forecast window, we employ the label Mean Squared Prediction Error (**MSPE**) before T_0 and the label **MSFE** from T_0 onward.

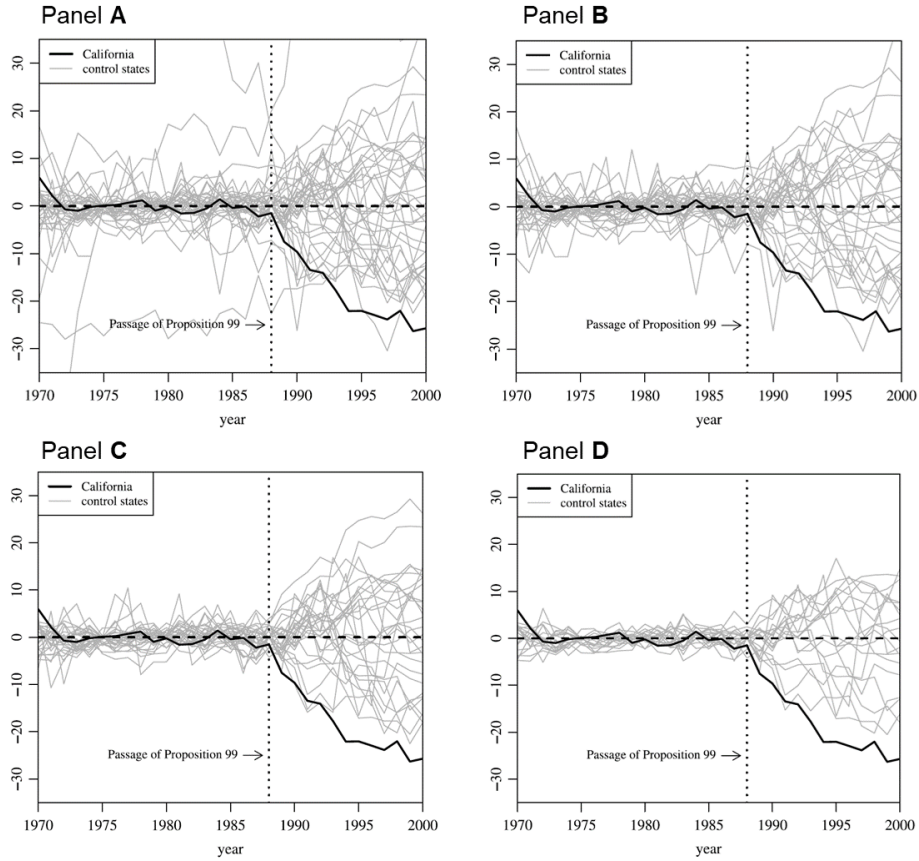


Figure 1. Region Exclusion Procedure of ADH

The vertical axis indicates the gap between observed and estimated per capita cigarette sales, the bold line represents the truly treated region (California). Two observations stand out when considering panel A: First of all, the treatment has a clear negative effect for California. Second, some regions have both a poor pre- and post-treatment fit. Since the treatment significance should not be artificially driven by regions with poor fit, **ADH** successively remove regions with a large **MSPE** relative to California. Panel B excludes regions with a **MSPE** that is more than 20 times as large the **MSPE** of California, Panel C lowers the cutoff to five times California's **MSPE** and Panel D to two times the **MSPE**. In the last scenario, only 19 regions are left and California is the one with the most extreme treatment effect. The authors therefore conclude that the treatment is statistically significant with a (permutation) p-value of 5,3% $\left(\frac{1}{19}\right)$.

One way to bypass the inefficient sample reduction procedure is to look at the distribution of the ratios of **MSFE** and **MSPE**. By scaling the post-treatment fit by the

pre-treatment fit, regions with a poor fit are implicitly controlled for. In the tobacco control application, California is the region with the highest **MSFE**-to-**MSPE** ratio among all 39 regions which translates into a p-value of 2,6% $\left(\frac{1}{39}\right)$.

3.2. Simple Static Extension

To give an intuitive introduction to our proposed extensions, we first consider the most simple scenario of one treatment unit $j = 0$ and two donor units $j = 1, 2$. We consider a setting where only the outcome series (e.g. **GDP**) and no further covariates (e.g. consumption, investment etc.) are observed. It is assumed that before $t = T_0$ the units have a joint distribution of the form

$$Y = \begin{pmatrix} y_0 \\ y_1 \\ y_2 \end{pmatrix} \sim \mathcal{N}(\mu, \Sigma) \text{ for } t < T_0.$$

with $\mu = (\mu_0, \mu_1, \mu_2)'$ and the positive definite covariance matrix

$$\Sigma = \begin{pmatrix} \sigma_0^2 & \sigma'_{12} \\ \sigma_{12} & \Sigma_2 \end{pmatrix}.$$

σ_0^2 denotes the variance of y_0 , Σ_2 is a (2×2) covariance matrix of the vector $(y_1, y_2)'$ and σ_{12} is a (2×1) vector with elements $cov(y_0, y_1)$ and $cov(y_0, y_2)$.

Disregarding any constraints, we are interested to derive the best unbiased forecast of y_0 given the controls y_1 and y_2 which is obtained as

$$\begin{aligned} \hat{y}_0^N &= \mu_0 + w_1^{OLS}(y_1 - \mu_1) + w_2^{OLS}(y_2 - \mu_2) \\ &= \mu^* + w_1^{OLS}y_1 + w_2^{OLS}y_2 \end{aligned}$$

where $\mu^* = \mu_0 - w_1^{OLS}\mu_1 - w_2^{OLS}\mu_2$. This forecast can be directly estimated by an unrestricted **OLS** regression of y_0 on y_1 and y_2 . However, the result implies that there is no inherent reason to impose the restrictions that $w_1^{OLS}, w_2^{OLS} \geq 0$ and $w_1^{OLS} + w_2^{OLS} = 1$. Furthermore, we argue that the construction of **SC** should include a constant term, as otherwise the estimated counterfactual may have a mean outside the convex hull of the

donor means. See also [Doudchenko and Imbens, 2016] for a careful discussion of these restrictions.

For illustrative reasons, assume that

$$Y \sim \mathcal{N} \left(\begin{pmatrix} 1 \\ 1 \\ 1 \end{pmatrix}, \begin{pmatrix} 1 & 0.1 & 0.4 \\ 0.1 & 1 & 0.5 \\ 0.4 & 0.5 & 1 \end{pmatrix} \right).$$

For this example the unrestricted optimal weights for the counterfactual result as $w_1^{OLS} = -0.1333$, $w_2^{OLS} = 0.4667$ and $\mu^* = \mu_0 - w_1^{OLS} \cdot \mu_1 - w_2^{OLS} \cdot \mu_2 = 0.6667$.⁷ Note that w_1^{OLS} is negative even though all bivariate correlations between the units are positive. One may argue that this result does not make much sense as the economic interpretation of y_1 entering the counterfactual \hat{y}_0^N with a negative sign is unclear. This demonstrates the trade-off between optimality in a statistical sense and the economic interpretation of the solution.

What happens if we impose the restrictions that all weights are positive and sum up to unity? In this case the restricted optimum yields the linear combination $\tilde{y}_0^N = 0.2y_1 + 0.8y_2$. The important difference lies in the variance of these estimates. For our example we obtain

$$var(y_0 - \hat{y}_0^N) = 0.8267$$

$$var(y_0 - \tilde{y}_0^N) = 1.1600.$$

It is interesting to note that the variance of the restricted estimate is even larger than the unconditional variance of y_0 . This is possible as $(w_1, w_2) = (0, 0)$ is not included in the restricted parameter space.

So far we argued and showed illustratively, that an unrestricted **OLS** estimate can be superior to the constraint **SC** estimate in settings with few panel units and a clear correlation structure among the units. This indication will be further refined in subsequent Monte Carlo simulations. In microeconomic settings it is usually assumed that the units in the treatment group and units in the control group are uncorrelated. In such cases the

⁷ The derivation of the employed estimators is postponed to the appendix.

construction of a **SC** is unpromising as the dependency between treatment unit and donors is the core condition for a plausible estimation of the counterfactual. In macroeconomic applications however, the variables in the treatment and control group (e.g. **GDP**) are typically correlated and it is therefore important to model this relationship. As the simple scenario with only two panel units in the donor pool is highly unrealistic in practice, we now move to the general static case with $J + 1$ panel units.

3.3. General Static Extension

In empirical macroeconomic practice, the observed time series are typically low- frequency, i. e. the quantities of interest are measured at monthly, quarterly or even annual intervals. Thus, the number of pre-intervention time periods ($T_0 - 1$) is typically small and may even be smaller than the number of units in the donor pool J . In such scenarios, the unrestricted **OLS** estimate may face issues of instability or, in the case of $T_0 - 1 < J$, due to singularity, it may not even be identified. The issues of external validity and overfitting are closely related to the aspect of identification. Especially when employing non-parametric statistical learning methods, it is simple to achieve a high in-sample (pre-treatment) fit. The crucial part when dealing with forecasts is that the observed in-sample patterns generalize well outside the verifiable horizon (post-treatment). **ADH** solve this issue by restricting the weights to be non-negative and to sum up to one. Besides preventing the model from overfitting, the percent restriction guarantees the existence of unique weights, especially when dealing with a small number of pre-treatment periods. Regularized regressions constitute another model family that is capable of balancing the trade-off between under- and overfitting.

ELASTIC NET

In this context [[Doudchenko and Imbens, 2016](#)] suggest employing an elastic net regression to regularize the donor weights. It solves the following objective function:

$$Q(w, \lambda_1, \lambda_2) = \sum_{t=1}^{T_0-1} \underbrace{\left(y_{0,t} - \mu^* - \sum_{j=1}^J w_j y_{j,t} \right)^2}_{RSS} + \lambda_1 \underbrace{\left(\sum_{j=1}^J w_j^2 \right)}_{Ridge} + \lambda_2 \underbrace{\left(\sum_{j=1}^J |w_j| \right)}_{Lasso}$$

The L_2 -norm (Ridge-Penalty) is a continuous shrinkage method, that shrinks the co-

efficients towards zero without performing variable selection in the sense that certain coefficients are set exactly to zero ([Hoerl and Kennard, 1970]). However, it has the appealing feature that its estimation only involves the addition of a diagonal matrix to the Residual Sum of Squares (**RSS**). Therefore, the objective function keeps an explicit closed form solution which is particularly important if the sample is small.

In contrast, the L_1 -norm (Lasso-Penalty) as proposed by [Tibshirani, 1996] penalizes the sum of the absolute values of the coefficients. The nature of the penalty term causes this regularization to perform both, continuous shrinkage and automatic variable selection. As a consequence, the argmax vector of the objective function typically contains many entries that are exactly zero which makes the resulting model sparse and easier to interpret. However, since the absolute value function is not continuously differentiable, the Lasso has no closed form solution. Consequently, the minimum of the objective function has to be approximated, which is typically done via numerical optimization techniques like cyclical coordinate descent algorithm (see for example [Friedman et al., 2010]). The shrinkage parameters λ_1 and λ_2 can be selected through k-fold cross-validation. This involves storing combinations of λ_1 and λ_2 that minimize the objective function across k validation sets. The average value of these hyperparameters is then computed to make the final choice.

REGSC

We propose a different regularization that we call the regularized synthetic control estimator. This estimator augments the **OLS** objective function by a Ridge penalty and a simple "inverse" Ridge that shrinks the coefficient sum towards one. The objective function has the following form:

$$Q(w, \lambda_1, \lambda_2) = \sum_{t=1}^{T_0-1} \left(y_{0,t} - \mu^* - \sum_{j=1}^J w_j y_{j,t} \right)^2 + \lambda_1 \left(\sum_{j=1}^J w_j^2 \right) + \lambda_2 \left(1 - \sum_{j=1}^J w_j \right)^2$$

Due to the individual shrinkage to zero and the joint shrinkage to one, this regularization is closely related to original **SC** estimator but in contrast to the elastic net, it is flexible enough to produce non-zero weights that are directly interpretable. Moreover, as it does not involve approximating the gradient of the absolute value function, it has the following

closed form solution:

$$\hat{w}_{\lambda_1, \lambda_2} = (Y_1' Y_1 + \lambda_1 I_J + \lambda_2 \mathbf{1}_J \mathbf{1}_J')^{-1} (Y_1' Y_0 + \lambda_2 \mathbf{1}_J)$$

I_J depicts a J -dimensional vector of ones, $\mathbf{1}_J \mathbf{1}_J'$ an all-ones matrix of dimension J . Consistent with the notation of chapter 3.1, Y_1 is a $(T_0 - 1) \times J$ matrix that stacks all pre-treatment donor observations for $t = 1, \dots, T_0 - 1$ and $j = 1, \dots, J$. Analogously, Y_0 is a $(T_0 - 1) \times 1$ vector that stacks the pre-treatment time series observations of the treatment unit. Note that Y_0 and Y_1 should be demeaned prior to the estimation to omit the constant from the penalty. Alternatively, a vector of ones should be joined to Y_1 from the left and the first element of I_J respectively $\mathbf{1}_J \mathbf{1}_J'$ be set to zero. Similar to the case of the elastic net, the shrinkage parameters λ_1 and λ_2 can be chosen by cross validation, where our experience suggests that optimizing subject to the restriction $\lambda_1 \approx 10000 \cdot \lambda_2$ reduces computation time and already produces reasonable estimates.

The combination of a closed form solution and tuneable hyperparameters make the **REGSC**-method highly appropriate in the low-frequency macroeconomic context of **SC**: It is able to produce weights that are interpretable, flexible and efficiently estimated in small samples. These characteristics are empirically verified in the subsequent simulation study. Besides the proposed regularization, we also implemented a numerically solvable combination of the Lasso- and the "inverse"-Ridge-penalty. In large data sets of at least 1,000 observations, this alternative was competitive to the elastic net and the proposed **REGSC**-estimator. However, as we are looking for an estimator with robust small sample properties, the Lasso-"inverse"-Ridge estimator is omitted from the further analysis.

3.4. General Dynamic Extension

When modeling macroeconomic time series it is often assumed that the $(J + 1) \times 1$ vector of time series $y_t = (Y_{0,t}, \dots, Y_{J,t})'$ can be represented by a **VAR** model of the following form:

TBD

4. Simulation

In this chapter, we empirically test the performance of our proposed and the already existing estimators for **SC** in different **DGP**. Independent of the specific features of the data in terms of pre- and post-treatment period length and the prevailing time series structures, we proceed as follows: We simulate $T_{pre} = T_0 - 1$ periods of pre-treatment and $T_{post} = T - (T_0 - 1)$ periods of post-treatment data for the single treated unit and the J donor units. Each estimator's main goal is to grasp the consistent patterns before treatment and accurately extend these patterns into the time after treatment. Said differently, the pre-treatment phase depicts the training set of the models and the post-treatment the validation set. To root the simulation framework as close as possible to real-world **SC** applications, we define T_{pre} and T_{post} such that their range is comparable to low-frequency macroeconomic settings, i.e. $T_{pre} \in \{20, 50, 100\}$ and $T_{post} \in \{10, 20, 30\}$. Furthermore, we consider two types of **DGP**, a static factor process and a dynamic **VAR**-process that is inspired by real **GDP** processes of the G20 countries.

4.1. Static Data Generating Processes

In their **SC** application of estimating the causal effect of California's proposition 99, [Abadie et al., 2010] suppose that the (potential) outcome $Y_{i,t}^N$ follows a factor model of the form

$$Y_{i,t}^N = \alpha_t + \theta_t Z_i + \lambda_t \mu_i + \epsilon_{it}.$$

α_t denotes an unknown panel-invariant factor, Z_i is a vector of observed panel-specific covariates, θ_t is a vector of unknown parameters, λ_t is a vector of unknown common factors and μ_i are panel-specific unknown factor loadings. The unobserved shocks ϵ_{it} have zero mean at the panel level. For this specific setting, [Abadie et al., 2010] show that "[...] the bias of the SC-estimator can be bounded by a function that goes to zero as the number of pre-treatment periods increases." Further the number of donor units has to be fixed. The fact, that α_t is panel-invariant seems minor at first glance. However, as the **SC**-estimator does neither contain an intercept nor does it allow for extrapolation outside the convex hull of the donor pool, the unbiasedness of the estimator directly depends on the distribution of the intercepts. In a slightly more realistic data-generating scenario, the

intercepts do not follow a degenerate point distribution with $P(X = 0) = 1$ but are drawn from a symmetric distribution centered around the origin like the standard normal.

[Ferman, 2021] considers a de-meaned scenario without additional covariates. In our simulation, we follow the basic set-up of Ferman and generate data according to a similar factor model. However, we consider it more realistic to add a time-invariant and panel-specific intercept to the (potential) outcome instead of analyzing a de-meaned **DGP**. Our representation of the counterfactuals therefore boils down to

$$Y_{i,t}^N = \alpha_i + \lambda_t \mu_i + \epsilon_{it}.$$

In this simplified setting, the counterfactual is given by the composition of the unknown panel-specific factor loadings μ_i and the F unknown common factors $\lambda_t = (\lambda_{1,t}, \dots, \lambda_{F,t})$ plus intercept α_i and idiosyncratic shocks ϵ_{it} . For the sake of simplicity, Ferman considers a scenario with only two common factors, $\lambda_{1,t}$ and $\lambda_{2,t}$. We proceed analogously and generates the data such that the (potential) outcome of the treated unit and the first half of the donor pool load exclusively with loading one on the first factor, the remaining donors load exclusively with loading one on the second factor. Therefore μ_i is a (2×1) -column-vector with the first (second) entry being one and the second (first) entry being zero for the first (second) half of the donor pool. Further, the random variables α_i , $\lambda_{1,t}$, $\lambda_{2,t}$ and ϵ_{it} are realizations of independent and identically distributed (**iid**) standard normal distributions $\mathcal{N}(0, 1)$. The following figure exemplifies the functionality of the **DGP** with $T_{pre} = 20$ and $T_{post} = 10$ and a constant treatment effect of $\delta_{1,t} = 10$ for $t > T_0$.⁸

⁸ In this example, the series $j = 0$ is treated at $T_0 = 20$, while the impact of the treatment becomes noticeable one time period later, starting from $t > T_0$. Note that the actual treatment effect is irrelevant for our investigation as it is empirically observable.




Figure 2. Example Factor-DGP

To make the factor structure easy to see, we scaled the factor variance by 10^1 and the error variance by 10^{-1} . The generated data exhibits a clearly observable factor structure: The treatment unit and the first half of the donors (Donor 1 and 2) as well as the second half of the donors (Donor 3 and 4) share a common factor. Thus, the objective of each employed method is to recover the true factor structure, i. e. irrelevant of the size of donor pool to weight only the first $\frac{J}{2}$ donors positively. Further, we see that each series possesses an own intercept. Yet, the intercept variation is dominated by the factor variation as in this example DGP only, there variance has been scaled by 10.

EMPLOYED MODELS

For the static factor DGP, we consider five models:

1. SC: The first model is the ordinary SC method without additional covariates. Therefore, this method is equivalent to a restricted OLS regression that regresses the treatment series on the donors series given the constraint of no intercept and non-negative coefficients that sum up to one. The belief for this model is that the untuneable restrictions prevent it from overfitting the pre-treatment data but that this inelasticity comes at the costs of a reduced predictive performance.
2. OLS: The second model is a usual regression that regresses the treatment series

on the donor series. The belief for this model is that it starts overfitting the pre-treatment data as J grows large. Further for $J > T_0$, it is unable to provide any prediction or forecast.

3. NET: The third model we consider is the aforementioned mentioned elastic net as proposed by [Doudchenko and Imbens, 2016] in the context of SC. Due to its flexible hyperparameters tuned via pre-treatment cross-validation, we expect this model to perform reasonable well both before and after the treatment. In the conceptual introduction of the elastic net, we stressed the potential drawback of having no closed form solution. As a consequence, we assume the model to perform worse in small samples.
4. REGSC: The fourth model under consideration is our proposed regularized synthetic control model which can be interpreted as a mixture of the elastic net and the original SC estimator. It is comparable to the elastic net insofar that it simply substitutes the Lasso-shrinkage by the inverse-Ridge-shrinkage. This regularization is motivated by the SC specificity of having percent-like coefficients. Thus, we expect the model to perform well especially in settings that are comparable to the original SC setting like the static factor DGP. Its closed form solution and the increased flexibility caused by the tuneable hyperparameters make us confident that the model performs equally well in small and in large as well as in the pre- and the post-treatment period.
5. FACTOR: The data of this chapter is generated such that the common factors and the idiosyncratic component are uncorrelated and that the idiosyncratic errors are mutually uncorrelated. Thus, the static factor model which estimates the common factors as linear function of the donors is the most natural candidate model. A popular estimator for the factor model is the Principal Components (PC) estimator. In the pre-treatment period, we obtain the predictions by regressing the treatment series on the "latent" factors. As we implemented a two-factor structure, the factors are computed by multiplying the first two eigenvalues of the covariates matrix with the matrix of the covariates. The forecasts for the post-treatment period are obtained by multiplying the factor structure of the post-treatment period with the regression coefficients of the pre-treatment regression. As this model directly build upon the

DGP, it is our "benchmark-model" and we expect it to perform best among all candidates. (why no overfitting?)

(what else?)

- All simulations are conducted with 1,000 iterations per combination of pre- and post-treatment horizon and donor quantity. (Important to distinguish different intercept cases: If the intercept of the treated series falls outside the donor-intercepts, SC exhibits a bias as it does not allow for a panel-specific constant.)
- employ RMSE as loss function. Quadratic loss, tremendously important in practice. Reasonable approximation to realistic loss structures and mathematically convenient ([Diebold, 2017])

4.2. Dynamic Data Generating Processes

5. Simulation

5.1. Static Data Generating Processes

Simulation-Procedure

5.2. Weakly Dynamic Data Generating Processes

5.3. Dynamic Data Generating Processes

5.3.1. Set up

In order to rigorously evaluate the performance of both, the here proposed and existing SC estimators, a key milestone would be to test the estimators in simulated datasets which mimic the real world. As many of the previous studies have focused on the economic development following a treatment, it stands to reason that real world changes in GDP could serve as the basis for a close-to-reality inspired DGP. To ensure that the DGP is based on a relatively uniform reference group which inhibits significant amounts of commonalities and correlation, the data basis for this simulation consists of all countries from the G20 as well as the European Union (EU).⁹ **Achtung: R Paket richtig zitieren** The dataset is subjected to two additional filters: firstly, only countries with at least 40 years worth of GDP data remain in the dataset, and secondly, the time series selected must be stationary. The latter is tested using the Augmented Dickey Fuller Test on a 10% significance level. The remaining 22 countries are the base from which the close-to-reality datasets are simulated using a VAR model.

Two different approaches, which will be referred to in the following as the *micro approach* and *macro approach*, are being considered. To evaluate the proposed models in terms of an increasing number of donors as in the static case, the results naturally vary with the underlying base, the data is simulated from. It depends critically on whether the simulation is performed with the equal number of donors as the models are estimated with (micro approach), or whether the simulation is performed with a larger number of donors such that the model estimation is carried out with subsets of the whole simulation (macro

⁹ The GDP data is sourced from the World Bank's World Development Indicators, which is directly accessible via the WDI-Package in R using the ticker 'NY.GDP.PCAP.KD.ZG' (GDP Per Capita Growth Rate).

approach). The micro approach is a more controlled and stylized design, in which the model estimation is performed on the full simulated dataset. Simulation and estimation are therefore performed using the same amount of J donors in each iteration, i.e. is therefore internally consistent. In this case, adding more donors means adding complexity to the simulation as well as the estimation of the model. The macro approach is a more realistic approach in the sense, that by increasing the number of donors, new information is added to the model estimation. Hence, in the macro approach the dataset is always simulated with J_{max} donors, while the model estimation is performed using J donors, such that for all estimation iterations $J_{max} > J$. This approach reflects the effects of adding new information in terms of donors to the model estimation while the simulated dataset remains structurally the same.

For simplicity, the lag order of the simulation VAR process is set to $p = 2$ and to ensure consistency with the static DGP the same T_{pre} and T_{post} values are used in simulation. However due to dimensionality constraints of the VAR model, the number of donors is reduced to $J \in \{2, 4, 6, 8\}$ and $J_{max} = 14$. Therefore, in each iteration $J + 1$ (micro), or $J_{max} + 1$ (macro) countries are drawn at random from the 22-country base. One of the countries is randomly assigned to be the treatment unit while the remaining ones function as (potential) donors.

5.3.2. Employed Models

For the Dynamic **DGP**, we employ the models of the static case as reference, as well as the following three models reflecting the time series structure of the data.

VAR => reference model

Unidyn:

Multidyn/ADL Model: Hier Md1 und MD2 nur kurz nennen, Multidyn 3 dann aber als Essenz mit den entsprechenden Vorteilen den anderen gegenüber. Dann in den Results wieder kurz darauf beziehen (Ergebnisse der anderen im Appendix?)

refer to the metrics in the static case => keep short here

5.3.3. Results

Short description of the model specific parameters used
divide by the two DGPs and also compare the development when J increases
Full Table in Appendix
Start with general findings then details regarding the models

Micro View

Macro View

6. Applications

We consider three leading examples:

6.1. The Economic Costs of Conflict

[[Abadie and Gardeazabal, 2003](#)]




Figure 3. GDP per Capita for the (synthetic) Basque

6.2. Estimating the Effect of California's Tobacco Control Program

[Abadie et al., 2010]




Figure 4. Cigarette Sales per Capita for the (synthetic) California

6.3. The Economic Cost of the 1990 German Reunification

[Abadie et al., 2015]




Figure 5. GDP per Capita for the (synthetic) West Germany

7. Conclusion

- Some concluding remarks and an outlook
- Keep short, around 1-2 pages
- Natural extension: case with explanatory variables
- We advocate for an interval instead of an point forecast. Therefore also for an interval estimate of the treatment effect. Report more measures of uncertainty than permutation p-values.

References

- [Abadie, 2021] Abadie, A. (2021). Using synthetic controls: Feasibility, data requirements, and methodological aspects. *Journal of Economic Literature*, 59(2):391–425.
- [Abadie et al., 2010] Abadie, A., Diamond, A., and Hainmueller, J. (2010). Synthetic control methods for comparative case studies: Estimating the effect of california’s tobacco control program. *Journal of the American Statistical Association*, 105:493–505.
- [Abadie et al., 2015] Abadie, A., Diamond, A., and Hainmueller, J. (2015). Comparative politics and the synthetic control method. *American Journal of Political Science*, 59(2):495–510.
- [Abadie and Gardeazabal, 2003] Abadie, A. and Gardeazabal, J. (2003). The economic costs of conflict: A case study of the basque country. *American Economic Review*, 93:113–132.
- [Abadie and Imbens, 2002] Abadie, A. and Imbens, G. (2002). Bias-corrected matching estimators for average treatment effects. *Journal of Business and Economic Statistics*, 29.
- [Abadie and Imbens, 2006] Abadie, A. and Imbens, G. (2006). Large sample properties of matching estimators for average treatment effects. *Econometrica*, 74:235–267.
- [Abadie and L’Hour, 2021] Abadie, A. and L’Hour, J. (2021). A penalized synthetic control estimator for disaggregated data. *Journal of the American Statistical Association*, 116:1–34.
- [Amjad et al., 2018] Amjad, M., Shah, D., and Shen, D. (2018). Robust synthetic control. *Journal of the American Statistical Association*, pages 1–51.
- [Andrews, 2003] Andrews, D. W. K. (2003). End-of-Sample Instability Tests. *Econometrica*, 71(6):1661–1694.
- [Arkhangelsky et al., 2021] Arkhangelsky, D., Athey, S., Hirshberg, D. A., Imbens, G. W., and Wager, S. (2021). Synthetic difference-in-differences. *American Economic Review*, 111(12):4088–4118.

- [Athey et al., 2017] Athey, S., Bayati, M., Doudchenko, N., Imbens, G., and Khosravi, K. (2017). Matrix completion methods for causal panel data models. *Journal of the American Statistical Association*, 116.
- [Athey and Imbens, 2016] Athey, S. and Imbens, G. (2016). The state of applied econometrics - causality and policy evaluation. *Journal of Economic Perspectives*, 31.
- [Ben-Michael et al., 2021] Ben-Michael, E., Feller, A., and Rothstein, J. (2021). The augmented synthetic control method. *SSRN Electronic Journal*.
- [Ben-Michael et al., 2021] Ben-Michael, E., Feller, A., and Rothstein, J. (2021). Synthetic controls with staggered adoption. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 84.
- [Born et al., 2019] Born, B., Müller, G. J., Schularick, M., and Sedláček, P. (2019). The Costs of Economic Nationalism: Evidence from the Brexit Experiment*. *The Economic Journal*, 129(623):2722–2744.
- [Breitung and Knüppel, 2021] Breitung, J. and Knüppel, M. (2021). How far can we forecast? Statistical tests of the predictive content. *Journal of Applied Econometrics*, 36(4):369–392.
- [Brodersen et al., 2015] Brodersen, K. H., Gallusser, F., Koehler, J., Remy, N., and Scott, S. L. (2015). Inferring causal impact using bayesian structural time-series models. *The Annals of Applied Statistics*, 9(1):247–274.
- [Cattaneo et al., 2021] Cattaneo, M., Feng, Y., and Titiunik, R. (2021). Prediction intervals for synthetic control methods*. *Journal of the American Statistical Association*, 116:1–44.
- [Chernozhukov et al., 2019] Chernozhukov, V., Wüthrich, K., and Zhu, Y. (2019). Inference on average treatment effects in aggregate panel data settings. CeMMAP working papers CWP32/19, Centre for Microdata Methods and Practice, Institute for Fiscal Studies.
- [Chernozhukov et al., 2021] Chernozhukov, V., Wüthrich, K., and Zhu, Y. (2021). An Exact and Robust Conformal Inference Method for Counterfactual and Synthetic Controls.

- University of California at San Diego, Economics Working Paper Series qt90m9d66s, Department of Economics, UC San Diego.
- [Cho, 2020] Cho, S.-W. S. (2020). Quantifying the impact of nonpharmaceutical interventions during the COVID-19 outbreak: The case of Sweden. *The Econometrics Journal*, 23(3):323–344.
- [Cunningham, 2021] Cunningham, S. (2021). *Causal Inference: The Mixtape*. Yale University Press.
- [Diebold, 2017] Diebold, F. X. (2017). *Forecasting*. Department of Economics, University of Pennsylvania.
- [Doudchenko and Imbens, 2016] Doudchenko, N. and Imbens, G. W. (2016). Balancing, Regression, Difference-In-Differences and Synthetic Control Methods: A Synthesis. NBER Working Papers 22791, National Bureau of Economic Research, Inc.
- [Ferman, 2021] Ferman, B. (2021). On the Properties of the Synthetic Control Estimator with Many Periods and Many Controls. *Journal of the American Statistical Association*, 116(536):1764–1772.
- [Firpo and Possebom, 2018] Firpo, S. and Possebom, V. (2018). Synthetic control method: Inference, sensitivity analysis and confidence sets. *Journal of Causal Inference*, 6(2).
- [Fisher, 1935] Fisher, R. A. (1971 [1935]). The design of experiments (9th ed.). *Macmillan*.
- [Frangakis and Rubin, 2002] Frangakis, C. E. and Rubin, D. B. (2002). Principal stratification in causal inference. *Biometrics*, 58(1):21–29.
- [Friedman et al., 2010] Friedman, J. H., Hastie, T., and Tibshirani, R. (2010). Regularization Paths for Generalized Linear Models via Coordinate Descent. *Journal of Statistical Software*, 33(i01).

- [Funke et al., 2020] Funke, M., Schularick, M., and Trebesch, C. (2020). Populist leaders and the economy. ECONtribute Discussion Papers Series 036, University of Bonn and University of Cologne, Germany.
- [Hahn and Shi, 2017] Hahn, J. and Shi, R. (2017). Synthetic control and inference. *Econometrics*, 5(4).
- [Hainmueller et al., 2011] Hainmueller, J., Diamond, A., and Abadie, A. (2011). Synth: An r package for synthetic control methods in comparative case studies. *Journal of Statistical Software*, 42.
- [Hartford et al., 2017] Hartford, J., Lewis, G., Leyton-Brown, K., and Taddy, M. (2017). Deep iv: A flexible approach for counterfactual prediction. In *Proceedings of the 34th International Conference on Machine Learning - Volume 70*, ICML’17, page 1414–1423. JMLR.org.
- [Harvey and Thiele, 2020] Harvey, A. and Thiele, S. (2020). Cointegration and control: Assessing the impact of events using time series data. *Journal of Applied Econometrics*, 36.
- [Hoerl and Kennard, 1970] Hoerl, A. E. and Kennard, R. W. (1970). Ridge regression: Biased estimation for nonorthogonal problems. *Technometrics*, 12:55–67.
- [Kellogg et al., 2021] Kellogg, M., Mogstad, M., Pouliot, G. A., and Torgovitsky, A. (2021). Combining matching and synthetic control to tradeoff biases from extrapolation and interpolation. *Journal of the American Statistical Association*, 116(536):1804–1816. PMID: 35706442.
- [Kuosmanen et al., 2021] Kuosmanen, T., Zhou, X., Eskelinen, J., and Malo, P. (2021). Design Flaw of the Synthetic Control Method. MPRA Paper 106328, University Library of Munich, Germany.
- [Martin et al., 2012] Martin, V., Hurn, S., and Harris, D. (2012). *Econometric Modelling with Time Series: Specification, Estimation and Testing*. Themes in Modern Econometrics. Cambridge University Press.

- [Muhlbach and Nielsen, 2019] Muhlbach, N. S. and Nielsen, M. S. (2019). Tree-based Synthetic Control Methods: Consequences of moving the US Embassy. Papers 1909.03968, arXiv.org.
- [Neyman, 1923] Neyman, J. (1923). On the application of probability theory to agricultural experiments. essay on principles. section 9. *Statistical Science* 5, 4:465–472.
- [Rosenbaum and Rubin, 1983] Rosenbaum, P. R. and Rubin, D. B. (1983). The central role of the propensity score in observational studies for causal effects. *Biometrika*, 70:41–55.
- [Rubin, 1974] Rubin, D. (1974). Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology*, 66(5):688–701.
- [Tibshirani, 1996] Tibshirani, R. (1996). Regression shrinkage and selection via the lasso. *Journal of the Royal Statistical Society (Series B)*, 58:267–288.
- [von Brzeski et al., 2015] von Brzeski, V., Taddy, M., and Draper, D. (2015). Causal inference in repeated observational studies: A case study of ebay product releases.

8. Appendix

8.1. Simple Static Extension

8.1.1. OLS Solution

In case the population covariance matrix is observable, the OLS-coefficients can be directly derived from it: $(w_1^{OLS}, w_2^{OLS}) = \mathbf{\Sigma}_2^{-1} \boldsymbol{\sigma}_{12}$

8.1.2. SC Solution

The restricted solution is can directly be derived from the covariance matrix. The first index in the square brackets indicates the row, the second the column position.

$$\begin{aligned} w_1^{SC} &= (\boldsymbol{\sigma}_{12}'[1] - \boldsymbol{\sigma}_{12}'[2] - \mathbf{\Sigma}_2[2, 1] + \mathbf{\Sigma}_2[1, 1]) / (\mathbf{\Sigma}_2[1, 1] + \mathbf{\Sigma}_2[2, 2] - 2 * \mathbf{\Sigma}_2[1, 2]) \\ &= (0.1 - 0.4 - 0.5 + 1) / (1 + 1 - 2 * 0.5) = 0.2 \end{aligned}$$

$$\begin{aligned} w_2^{SC} &= (\boldsymbol{\sigma}_{12}'[2] - \boldsymbol{\sigma}_{12}'[1] - \mathbf{\Sigma}_2[1, 2] + \mathbf{\Sigma}_2[2, 2]) / (\mathbf{\Sigma}_2[2, 2] + \mathbf{\Sigma}_2[1, 1] - 2 * \mathbf{\Sigma}_2[2, 1]) \\ &= (0.4 - 0.1 - 0.5 + 1) / (1 + 1 - 2 * 0.5) = 0.8 \end{aligned}$$

8.1.3. Variances

The variances are derived from the weights and the covariance matrix:

$$\begin{aligned} var(Y_0 - w_1 Y_1 - w_2 Y_2) &= var(Y_0) + w_1^2 \cdot var(Y_1) + w_2^2 \cdot var(Y_2) - \\ &\quad 2 \cdot w_1 \cdot cov(Y_0, Y_1) - 2 \cdot w_2 \cdot cov(Y_0, Y_2) + \\ &\quad 2 \cdot w_1 w_2 \cdot cov(Y_1, Y_2) \end{aligned}$$

8.2. Simulation Study

8.2.1. Static Simulation results

Table S1. Simulation Results of the Static Factor Model with $\mathbf{J} = \mathbf{5}$ Donors.

T_0	T_1	FACTOR	SC	REGOLS	NET	OLS
		RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]
20	30	1.2494 (0.0166) [0.6158]	1.4604 (0.0001) [1.1078]	1.2942 (0.0142) [0.6464]	1.2959 (0.0123) [0.5695]	1.3552 (0.0102) [1.1093]
		1.2595 (-0.0068) [0.6477]	1.4560 (0.0320) [1.0652]	1.3080 (-0.0009) [0.6610]	1.3067 (-0.0003) [0.5948]	1.3665 (0.0066) [1.0979]
		1.2342 (0.0250) [0.6312]	1.4322 (0.0256) [1.0473]	1.2877 (0.0256) [0.6361]	1.2865 (0.0202) [0.5516]	1.3569 (0.0309) [1.0745]
	20	1.1819 (-0.0025) [0.6326]	1.4327 (-0.0015) [1.0579]	1.2130 (-0.0036) [0.6444]	1.1987 (-0.0014) [0.5430]	1.2133 (-0.0034) [0.7755]
		1.1837 (0.0092) [0.6469]	1.4224 (0.0097) [1.0563]	1.2124 (0.0144) [0.6396]	1.2046 (0.0131) [0.5583]	1.2194 (0.0150) [0.7840]
		1.1663 (0.0038) [0.5825]	1.3990 (-0.0019) [0.9754]	1.1966 (-0.0041) [0.5731]	1.1815 (0.0045) [0.5071]	1.1968 (-0.0014) [0.7163]
	10	1.1636 (-0.0046) [0.6505]	1.3944 (0.0176) [1.0736]	1.1818 (-0.0060) [0.6336]	1.1746 (-0.0052) [0.5802]	1.1807 (-0.0057) [0.7192]
		1.1684 (-0.0009) [0.6201]	1.3973 (0.0049) [1.0416]	1.1861 (-0.0024) [0.5941]	1.1808 (-0.0020) [0.5488]	1.1862 (-0.0022) [0.6838]
		1.1434 (0.0127) [0.5850]	1.3686 (0.0263) [0.9784]	1.1552 (0.0114) [0.5811]	1.1528 (0.0123) [0.5151]	1.1557 (0.0129) [0.6471]

Table S2. Simulation Results of the Static Factor Model with $\mathbf{J} = \mathbf{10}$ Donors.

T_0	T_1	FACTOR	SC	REGOLS	NET	OLS
		RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]
20	30	1.1613 (-0.0051) [0.7638]	1.2867 (-0.0034) [1.0520]	1.2304 (-0.0040) [0.7474]	1.2739 (-0.0057) [0.7799]	1.6198 (-0.0090) [2.1449]
		1.1555 (-0.0059) [0.7316]	1.2840 (-0.0267) [1.0055]	1.2250 (-0.0059) [0.7538]	1.2595 (0.0061) [0.7655]	1.6167 (-0.0024) [2.1752]
		1.1312 (-0.0091) [0.6970]	1.2532 (-0.0102) [0.9745]	1.1987 (-0.0140) [0.7263]	1.2395 (-0.0149) [0.7504]	1.5811 (-0.0052) [2.0458]
	20	1.1043 (0.0071) [0.8100]	1.2497 (0.0032) [1.0271]	1.1369 (0.0109) [0.7619]	1.1467 (0.0111) [0.6834]	1.2086 (0.0083) [1.1120]
		1.0966 (0.0079) [0.7854]	1.2308 (-0.0060) [1.0024]	1.1317 (0.0024) [0.7350]	1.1407 (0.0063) [0.6546]	1.2074 (0.0048) [1.0766]
		1.0912 (-0.0076) [0.7318]	1.2167 (-0.0053) [0.9551]	1.1204 (-0.0075) [0.7016]	1.1281 (-0.0058) [0.6186]	1.1895 (0.0001) [1.0395]
	10	1.0851 (0.0065) [0.7875]	1.2316 (0.0110) [1.0178]	1.1063 (0.0067) [0.7389]	1.1112 (0.0080) [0.6770]	1.1323 (0.0081) [0.9248]
		1.0733 (-0.0137) [0.7698]	1.2088 (0.0019) [0.9747]	1.0917 (-0.0134) [0.7290]	1.0941 (-0.0125) [0.6631]	1.1202 (-0.0122) [0.9043]
		1.0661 (-0.0082) [0.7524]	1.1986 (-0.0119) [0.9553]	1.0895 (-0.0096) [0.7098]	1.0934 (-0.0129) [0.6451]	1.1178 (-0.0113) [0.8864]

Table S3. Simulation Results of the Static Factor Model with $\mathbf{J} = \mathbf{15}$ Donors.

T_0	T_1	FACTOR	SC	REGOLS	NET	OLS
		RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]
20	30	1.1355 (0.0015) [0.7829]	1.2504 (-0.0022) [1.0159]	1.2171 (0.0004) [0.7942]	1.3089 (0.0012) [1.0400]	2.4509 (0.0009) [6.0410]
		1.1223 (0.0082) [0.7681]	1.2368 (-0.0057) [1.0051]	1.2088 (0.0130) [0.8111]	1.3010 (0.0024) [1.1466]	2.4339 (0.0089) [6.0128]
		1.1199 (0.0060) [0.7029]	1.2299 (0.0152) [0.9262]	1.2142 (-0.0096) [0.7565]	1.3095 (-0.0061) [1.0792]	2.4066 (-0.0104) [5.7301]
	20	1.0834 (0.0062) [0.8168]	1.2104 (0.0041) [0.9866]	1.1219 (0.0041) [0.7623]	1.1337 (0.0041) [0.6970]	1.2756 (0.0005) [1.3377]
		1.0794 (0.0047) [0.7768]	1.2016 (0.0050) [0.9709]	1.1222 (0.0043) [0.7428]	1.1324 (0.0054) [0.6599]	1.2737 (0.0025) [1.2865]
		1.0674 (0.0001) [0.7333]	1.1873 (-0.0077) [0.8756]	1.1010 (0.0021) [0.6879]	1.1105 (0.0026) [0.6355]	1.2532 (0.0028) [1.2290]
	10	1.0732 (0.0036) [0.8262]	1.1780 (0.0093) [0.9729]	1.0966 (0.0043) [0.7689]	1.1032 (0.0062) [0.6935]	1.1508 (0.0059) [1.0253]
		1.0645 (-0.0153) [0.8304]	1.1774 (-0.0091) [0.9807]	1.0905 (-0.0185) [0.7736]	1.0963 (-0.0192) [0.7014]	1.1459 (-0.0203) [1.0300]
		1.0627 (0.0114) [0.7628]	1.1613 (0.0090) [0.9068]	1.0829 (0.0146) [0.7123]	1.0902 (0.0149) [0.6393]	1.1378 (0.0171) [0.9584]
50	30	1.0732 (0.0036) [0.8262]	1.1780 (0.0093) [0.9729]	1.0966 (0.0043) [0.7689]	1.1032 (0.0062) [0.6935]	1.1508 (0.0059) [1.0253]
		1.0645 (-0.0153) [0.8304]	1.1774 (-0.0091) [0.9807]	1.0905 (-0.0185) [0.7736]	1.0963 (-0.0192) [0.7014]	1.1459 (-0.0203) [1.0300]
		1.0627 (0.0114) [0.7628]	1.1613 (0.0090) [0.9068]	1.0829 (0.0146) [0.7123]	1.0902 (0.0149) [0.6393]	1.1378 (0.0171) [0.9584]
	20	1.0732 (0.0036) [0.8262]	1.1780 (0.0093) [0.9729]	1.0966 (0.0043) [0.7689]	1.1032 (0.0062) [0.6935]	1.1508 (0.0059) [1.0253]
		1.0645 (-0.0153) [0.8304]	1.1774 (-0.0091) [0.9807]	1.0905 (-0.0185) [0.7736]	1.0963 (-0.0192) [0.7014]	1.1459 (-0.0203) [1.0300]
		1.0627 (0.0114) [0.7628]	1.1613 (0.0090) [0.9068]	1.0829 (0.0146) [0.7123]	1.0902 (0.0149) [0.6393]	1.1378 (0.0171) [0.9584]
	10	1.0732 (0.0036) [0.8262]	1.1780 (0.0093) [0.9729]	1.0966 (0.0043) [0.7689]	1.1032 (0.0062) [0.6935]	1.1508 (0.0059) [1.0253]
		1.0645 (-0.0153) [0.8304]	1.1774 (-0.0091) [0.9807]	1.0905 (-0.0185) [0.7736]	1.0963 (-0.0192) [0.7014]	1.1459 (-0.0203) [1.0300]
		1.0627 (0.0114) [0.7628]	1.1613 (0.0090) [0.9068]	1.0829 (0.0146) [0.7123]	1.0902 (0.0149) [0.6393]	1.1378 (0.0171) [0.9584]

Table S4. Simulation Results of the Static Factor Model with $\mathbf{J} = \mathbf{20}$ Donors.

T_0	T_1	FACTOR	SC	REGOLS	NET	OLS
		RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]
20	30	1.1278 (0.0020) [0.7964]	1.2332 (-0.0040) [1.0193]	1.2183 (-0.0039) [0.8413]	1.2859 (-0.0079) [0.9901]	NA NA NA
		1.1244 (-0.0082) [0.7614]	1.2312 (0.0043) [0.9685]	1.2083 (-0.0054) [0.7839]	1.2712 (0.0025) [0.9067]	NA NA NA
		1.1054 (0.0172) [0.7525]	1.1986 (0.0063) [0.9464]	1.1987 (0.0171) [0.8319]	1.2632 (0.0169) [0.9309]	NA NA NA
	20	1.0703 (-0.0056) [0.8408]	1.1749 (-0.0160) [0.9813]	1.1087 (-0.0070) [0.7851]	1.1247 (-0.0088) [0.7387]	1.3707 (-0.0067) [1.6546]
		1.0678 (0.0207) [0.8154]	1.1720 (0.0202) [0.9649]	1.1117 (0.0149) [0.7801]	1.1237 (0.0194) [0.7150]	1.3646 (0.0054) [1.5848]
		1.0514 (-0.0141) [0.7774]	1.1444 (-0.0264) [0.8997]	1.0846 (-0.0120) [0.7324]	1.0976 (-0.0085) [0.6672]	1.3227 (-0.0089) [1.4695]
	10	1.0447 (-0.0039) [0.8600]	1.1422 (0.0168) [0.9745]	1.0679 (-0.0029) [0.7904]	1.0793 (-0.0007) [0.7231]	1.1626 (-0.0016) [1.1338]
		1.0510 (-0.0144) [0.8576]	1.1476 (-0.0185) [0.9550]	1.0715 (-0.0139) [0.7864]	1.0815 (-0.0105) [0.7234]	1.1589 (-0.0113) [1.1391]
		1.0304 (0.0223) [0.8118]	1.1144 (0.0166) [0.9018]	1.0505 (0.0271) [0.7425]	1.0598 (0.0270) [0.6826]	1.1397 (0.0321) [1.0761]

Table S5. Simulation Results of the Static Factor Model with $\mathbf{J} = \mathbf{25}$ Donors.

T_0	T_1	FACTOR	SC	REGOLS	NET	OLS
		RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]
20	30	1.1093	NA	1.2010	1.2554	NA
		(0.0013)	NA	(0.0090)	(0.0116)	NA
		[0.7989]	NA	[0.8648]	[0.9421]	NA
	20	1.1073	NA	1.1989	1.2708	NA
		(0.0131)	NA	(0.0108)	(0.0093)	NA
		[0.8177]	NA	[0.8516]	[0.9724]	NA
	10	1.0971	NA	1.1928	1.2447	NA
		(-0.0046)	NA	(-0.0001)	(0.0013)	NA
		[0.7589]	NA	[0.8153]	[0.8945]	NA
50	30	1.0597	1.1582	1.0996	1.1172	1.4863
		(-0.0018)	(0.0135)	(-0.0048)	(-0.0079)	(-0.0151)
		[0.8466]	[0.9590]	[0.7879]	[0.7353]	[1.9825]
	20	1.0538	1.1525	1.0993	1.1139	1.4811
		(-0.0150)	(0.0043)	(-0.0082)	(-0.0101)	(0.0033)
		[0.8350]	[0.9332]	[0.7821]	[0.7263]	[1.9856]
	10	1.0432	1.1428	1.0831	1.0997	1.4400
		(0.0062)	(-0.0038)	(0.0062)	(0.0064)	(-0.0066)
		[0.8077]	[0.9111]	[0.7625]	[0.7214]	[1.8603]
100	30	1.0378	1.1232	1.0641	1.0741	1.1959
		(0.0000)	(0.0058)	(-0.0003)	(0.0006)	(0.0044)
		[0.8501]	[0.9454]	[0.7912]	[0.7149]	[1.2244]
	20	1.0483	1.1401	1.0747	1.0842	1.2002
		(0.0110)	(0.0130)	(0.0114)	(0.0127)	(0.0147)
		[0.8655]	[0.9457]	[0.7968]	[0.7234]	[1.2330]
	10	1.0296	1.1154	1.0551	1.0648	1.1811
		(-0.0018)	(-0.0017)	(-0.0012)	(0.0001)	(-0.0040)
		[0.8083]	[0.8879]	[0.7444]	[0.6782]	[1.1463]

Table S6. Simulation Results of the Static Factor Model with $\mathbf{J} = \mathbf{30}$ Donors.

T_0	T_1	FACTOR	SC	REGOLS	NET	OLS
		RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]	RMSE (BIAS) [VARIANCE]
20	30	1.1013	NA	1.1899	1.2523	NA
		(0.0128)	NA	(0.0149)	(0.0115)	NA
		[0.8224]	NA	[0.8865]	[0.9715]	NA
	20	1.1043	NA	1.1825	1.2442	NA
		(0.0014)	NA	(0.0031)	(0.0052)	NA
		[0.8072]	NA	[0.8431]	[0.9420]	NA
	10	1.0890	NA	1.1647	1.2337	NA
		(-0.0075)	NA	(-0.0054)	(-0.0055)	NA
		[0.7761]	NA	[0.8107]	[0.8943]	NA
50	30	1.0622	1.1477	1.1039	1.1297	1.6851
		(-0.0022)	(0.0062)	(0.0009)	(-0.0012)	(0.0059)
		[0.8658]	[0.9630]	[0.8137]	[0.7852]	[2.6638]
	20	1.0508	1.1372	1.0936	1.1190	1.6716
		(-0.0016)	(-0.0084)	(-0.0009)	(-0.0004)	(-0.0006)
		[0.8700]	[0.9523]	[0.8184]	[0.7841]	[2.5964]
	10	1.0445	1.1342	1.0880	1.1061	1.6681
		(0.0006)	(-0.0043)	(0.0052)	(0.0091)	(0.0174)
		[0.8279]	[0.9218]	[0.7784]	[0.7648]	[2.5590]
100	30	1.0420	1.1178	1.0662	1.0793	1.2394
		(0.0097)	(0.0172)	(0.0107)	(0.0094)	(0.0070)
		[0.8864]	[0.9693]	[0.8258]	[0.7565]	[1.3694]
	20	1.0435	1.1237	1.0669	1.0788	1.2351
		(0.0072)	(0.0038)	(0.0054)	(0.0057)	(0.0064)
		[0.8778]	[0.9520]	[0.8077]	[0.7417]	[1.3400]
	10	1.0159	1.0890	1.0410	1.0465	1.2026
		(0.0027)	(0.0173)	(0.0041)	(0.0020)	(0.0038)
		[0.8381]	[0.8876]	[0.7678]	[0.7076]	[1.2823]